


Step-by-Step Guide to Building a RAG Chatbot with Amazon Bedrock

Step 1: Create an S3 Bucket and Upload Documents

1. Go to the Amazon S3 console.
2. Click Create bucket.
 - Bucket name: nextwork-rag-bedrock-<your-initials>
 - Region: us-east-2
 - Leave default settings (ACLs disabled, block public access, default encryption).
3. Click Create bucket.
4. Open the bucket and click Upload.
5. Select your PDF or document files and upload.



aws

Search

[Alt+S]

United States (Ohio)

Ritesh Meshram

Storage

Amazon S3

Store and retrieve any amount of data from anywhere

Amazon S3 is an object storage service that offers industry-leading scalability, data availability, security, and performance.

Create a bucket

Every object in S3 is stored in a bucket. To upload files and folders to S3, you'll need to create a bucket where the objects will be stored.

Create bucket

How it works



Pricing

With S3, there are no minimum fees. You only pay for what you use. Prices are based on the location of your S3 bucket.

Estimate your monthly bill using the [AWS Simple Monthly Calculator](#)

[View pricing details](#)

Resources

aws

Search

[Alt+S]

United States (Ohio)

Ritesh Meshram

Amazon S3

Buckets

Create bucket

Create bucket [Info](#)

Buckets are containers for data stored in S3.

General configuration

AWS Region
US East (Ohio) us-east-2

Bucket type [Info](#)

☒ **General purpose**
Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ **Directory**
Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name [Info](#)

myawsbucket

Bucket names must be 3 to 63 characters and unique within the global namespace. Bucket names must also begin and end with a letter or number. Valid characters are a-z, 0-9, periods (.), and hyphens (-). [Learn More](#)

Copy settings from existing bucket - optional
Only the bucket settings in the following configuration are copied.

[Choose bucket](#)

Format: s3://bucket/prefix

Object Ownership [Info](#)

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

Search

[Alt+S]

United States (Ohio)Ritesh Meshram

Amazon S3BucketsCreate bucket

Bucket typeInfo

☒ General purpose

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ Directory

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket nameInfo

AWS-rag-S3-bucket

Bucket names must be 3 to 63 characters and unique within the global namespace. Bucket names must also begin and end with a letter or number. Valid characters are a-z, 0-9, periods (.), and hyphens (-). [Learn More](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

Choose bucket

Format: s3://bucket/prefix

Object OwnershipInfo

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

☒ ACLs disabled (recommended)

All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

☐ ACLs enabled

Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

Object Ownership

Bucket owner enforced

Block Public Access settings for this bucket

Search

[Alt+S]

United States (Ohio)Ritesh Meshram

Amazon S3BucketsCreate bucket

Object OwnershipInfo

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

☒ ACLs disabled (recommended)

All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

☐ ACLs enabled

Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

Object Ownership

Bucket owner enforced

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

☒ Block all public access

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☒ Block public access to buckets and objects granted through new access control lists (ACLs)

S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change any existing permissions that allow public access to S3 resources using ACLs.

☒ Block public access to buckets and objects granted through any access control lists (ACLs)

S3 will ignore all ACLs that grant public access to buckets and objects.

☒ Block public access to buckets and objects granted through new public bucket or access point policies

S3 will block new bucket and access point policies that grant public access to buckets and objects. This setting doesn't change any existing policies that allow public access to S3 resources.

☒ Block public and cross-account access to buckets and objects through any public bucket or access point policies

Search

[Alt+S]

United States (Ohio)Ritesh Meshram

Amazon S3BucketsCreate bucket

Default encryptionInfo

Server-side encryption is automatically applied to new objects stored in this bucket.

Encryption typeInfo

☒ Server-side encryption with Amazon S3 managed keys (SSE-S3)

☐ Server-side encryption with AWS Key Management Service keys (SSE-KMS)

☐ Dual-layer server-side encryption with AWS Key Management Service keys (DSSE-KMS)

Secure your objects with two separate layers of encryption. For details on pricing, see [DSSE-KMS pricing](#) on the Storage tab of the [Amazon S3 pricing page](#).

Bucket Key

Using an S3 Bucket Key for SSE-KMS reduces encryption costs by lowering calls to AWS KMS. S3 Bucket Keys aren't supported for DSSE-KMS. [Learn more](#)

☐ Disable

☒ Enable

Advanced settings

After creating the bucket, you can upload files and folders to the bucket, and configure additional bucket settings.

CancelCreate bucket

aws

Search

[Alt+S]

United States (Ohio)

Ritesh Meshram

Amazon S3

Buckets

Successfully created bucket "aws-rag-s3-bucket"

To upload files and folders, or to configure additional bucket settings, choose [View details](#).

Account snapshot - updated every 24 hours

All AWS Regions

View Storage Lens dashboard

General purpose buckets

Directory buckets

General purpose buckets (1)

Info

All AWS Regions

Buckets are containers for data stored in S3.

Find buckets by name

Copy ARN

Empty

Delete

Create bucket

1

Name

▲

AWS Region

▼

IAM Access Analyzer

Creation date

▼

aws-rag-s3-bucket

US East (Ohio) us-east-2

[View analyzer for us-east-2](#)

March 29, 2025, 22:19:34 (UTC+09:00)

aws

Search

[Alt+S]

United States (Ohio)

Ritesh Meshram

Amazon S3

Buckets

aws-rag-s3-bucket

aws-rag-s3-bucket

Info

Objects

Metadata

Properties

Permissions

Metrics

Management

Access Points

Objects (0)

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

1

Name

▲

Type

▼

Last modified

▼

Size

▼

Storage class

▼

No objects

You don't have any objects in this bucket.

Upload

aws

Search

[Alt+S]

United States (Ohio)

Ritesh Meshram

Upload succeeded

For more information, see the [Files and folders](#) table.

After you navigate away from this page, the following information is no longer available.

Summary

Destination

s3://aws-rag-s3-bucket

Succeeded

2 files, 41.3 MB (100.00%)

Failed

0 files, 0 B (0%)

Files and folders

Configuration

Files and folders (2 total, 41.3 MB)

Find by name

1

Name

Folder

▼

Type

▼

Size

▼

Status

▼

Error

▼

[AWS Security Cookbook.pdf](#)

-

application/pdf

23.6 MB

Succeeded

-

[Cloud Security Practical Guide t...](#)

-

application/pdf


17.7 MB

Succeeded

-

Step 2: Create a Knowledge Base in Bedrock

1. Open the Amazon Bedrock console.
2. In the left sidebar, click Knowledge Bases under Builder Tools.
3. Click Create, then select Knowledge Base with vector store.
4. Enter the following:
 - Knowledge base name: nextwork-rag-documentation
 - Description: This Knowledge Base stores all documentation at NextWork.
5. Under IAM permissions, select Create and use a new service role.
6. Under Choose data source, select Amazon S3.
7. Browse and select your S3 bucket (e.g., nextwork-rag-bedrock-<your-initials>).
8. Leave default parsing and chunking strategy.
9. Click Next.



Amazon Bedrock

<

▼ Getting started

Overview

Providers

▼ Foundation models

Model catalog [New](#)

Marketplace deployments [New](#)

▼ Playgrounds

Chat / Text

▼ Builder tools

Agents

Flows

Knowledge Bases

Prompt Management

▼ Safeguards

Guardrails

▼ Inference and Assessment

Batch inference

Cross-region inference

Evaluations

Amazon Bedrock > Knowledge Bases

Knowledge Bases

▼ How it works

Create a Knowledge Base

Knowledge Base with vector store: Build a fully customizable Knowledge Base with maximum flexibility. Specify the location of your data, select an embedding model, and configure a vector store. Bedrock stores and updates your embeddings.

Knowledge Base with structured data store: Build a Knowledge Base which can connect to a structured data source.

Create ▼

Test the Knowledge Base

Query your Knowledge Base in the test window. You can get source text chunks, or you can use the chunks to get responses from a foundation model.

Use the Knowledge Base

Integrate your Knowledge Base into your application as is or add it to agents.

Knowledge Bases

Edit Delete Test Knowledge Base Evaluate Create ▼

Find Knowledge Base

< 1 > ⚙

Name ▼

Status ▼

Type ▼

Data sources ▼

Description ▼

Creation time ▼

Last sync warn...

No Knowledge Base

Amazon Bedrock

<

▼ Getting started

Overview

Providers

▼ Foundation models

Model catalog [New](#)

Marketplace deployments [New](#)

▼ Playgrounds

Chat / Text

▼ Builder tools

Agents

Flows

Knowledge Bases

Prompt Management

▼ Safeguards

Guardrails

▼ Inference and Assessment

Batch inference

Cross-region inference

Evaluations

Amazon Bedrock > Knowledge Bases > Create knowledge base with vector store

Step 1 Provide Knowledge Base details

Step 2 Configure data source

Step 3 Configure data storage and processing

Step 4 Review and create

Provide Knowledge Base details

Knowledge Base details

Knowledge Base name

knowledge-base-quick-start-0ip6t

Valid characters are a-z, A-Z, 0-9, _ (underscore) and - (hyphen). The name can have up to 50 characters.

Knowledge Base description - optional

Enter description

Valid characters are a-z, A-Z, 0-9, _ (underscore) and - (hyphen). The name can have up to 200 characters.

IAM permissions

Certain permissions are necessary to access other services or perform actions in order to create this resource. For more information, see [service role](#) for Amazon Bedrock.



Runtime role

☒ Create and use a new service role

☐ Use an existing service role

Service role name

AmazonBedrockFoundationRoleForKnowledgeBase-0ip6t



Amazon Bedrock

- Getting started
 - Overview
 - Providers
- Foundation models
 - Model catalog [New](#)
 - Marketplace deployments [New](#)
- Playgrounds
 - Chat / Text
- Builder tools
 - Agents
 - Flows
 - Knowledge Bases
 - Prompt Management
- Safeguards
 - Guardrails
- Inference and Assessment
 - Batch inference
 - Cross-region inference
 - Evaluations

S3 URI

To increase the accuracy and relevance of your responses, add a metadata.json file containing metadata for your data source to your S3 bucket. [Info](#)

[View](#) [Browse S3](#)

☐ Add customer-managed KMS key for S3 data - *optional*
If you encrypted your S3 data, provide the KMS key here so that Bedrock can decrypt it.

Parsing strategy
Select a parsing option to configure how your data is processed. You can't modify this option after the knowledge base is created.

☒ **Amazon Bedrock default parser**
Select this option to only process text from your data. This parser doesn't incur charges.

☐ **Foundation models as a parser**
Select this option to process visually rich documents or images. With this option, you can use the default parser prompt or customize it for your use case. [Pricing details](#)

Chunking strategy
Chunking breaks down the text into smaller segments before embedding. The chunking strategy can't be modified after you create the data source.

Automatically splits text into chunks of about 300 tokens in size, by default. If a document is less than or already 300 tokens, it's not split...

► **Transformation function - optional**

► **Advanced settings - optional**

[Add data source](#)
You can add 4 more data source(s).

[Cancel](#) [Previous](#) [Next](#)

Step 3: Set Embeddings Model and Vector Store

1. In the Embeddings model section, click Select model.
2. Choose Titan Text Embeddings v2 and click Apply.
3. Under Vector store creation method, select Quick create a new vector store.
4. Under Vector store, choose Amazon OpenSearch Serverless.
5. Click Next and review your Knowledge Base setup.
6. Click Create Knowledge Base.

Amazon Bedrock

- Getting started
 - Overview
 - Providers
- Foundation models
 - Model catalog [New](#)
 - Marketplace deployments [New](#)
- Playgrounds
 - Chat / Text
- Builder tools
 - Agents
 - Flows
 - Knowledge Bases
 - Prompt Management
- Safeguards
 - Guardrails
- Inference and Assessment
 - Batch inference
 - Cross-region inference
 - Evaluations

Amazon Bedrock > Knowledge Bases > Create knowledge base with vector store

Step 1: Provide Knowledge Base details
Step 2: Configure data source
Step 3: Configure data storage and processing
Step 4: **Review and create**

Review and create

Step 1: Provide details [Edit](#)

Knowledge Base details

Knowledge Base name AWS-rag-documentation	Knowledge Base description —	Service role AmazonBedrockExecutionRoleForKnowledgeBase_rueva
Knowledge base type Knowledge base use vector store	Data source type Amazon S3	Log Deliveries —

Step 2: Configure data source [Edit](#)

Data source: AWS-rag-s3-bucket

Data source name AWS-rag-s3-bucket	Account ID 593793030302 (this account)	S3 URI s3://aws-rag-s3-bucket
--	--	---

[Previous](#) [Next](#)

Amazon Bedrock

▼ Getting started

Overview

Providers

▼ Foundation models

Model catalog [New](#)

Marketplace deployments [New](#)

▼ Playgrounds

Chat / Text

▼ Builder tools

Agents

Flows

Knowledge Bases

Prompt Management

▼ Safeguards

Guardrails

▼ Inference and Assessment

Batch inference

Cross-region inference

Evaluations

DELETE

Step 3: Select embeddings model and configure vector store

Embeddings model

Model	Embedding type	Vector dimensions
Titan Text Embeddings v2	Float vector embeddings	1024

Vector store

Quick create vector store - *Recommended*

Amazon OpenSearch Serverless

Multimodal storage destination

S3 URI

-

CancelPreviousCreate Knowledge Base

Amazon Bedrock

▼ Getting started

Overview

Providers

▼ Foundation models

Model catalog [New](#)

Marketplace deployments [New](#)

▼ Playgrounds

Chat / Text

▼ Builder tools

Agents

Flows

Knowledge Bases

Prompt Management

▼ Safeguards

Guardrails

[Amazon Bedrock](#) > [Knowledge Bases](#) > AWS-rag-documentation


AWS-rag-documentation

TestDelete

Knowledge Base overview

Edit

Knowledge Base name AWS-rag-documentation	Knowledge Base ID 2UXNKTNUUY
Knowledge Base description this knowledge base is based on documentation provided on AWS security	Status Available
Service Role AmazonBedrockExecutionRoleForKnowledgeBase_297u7	Created date March 30, 2025, 07:36 (UTC+09:00)
Log Deliveries Configure log deliveries and event logs in the Edit page.	
Retrieval-Augmented Generation (RAG) type Vector store	



Step 4: Enable Access to AI Models in Bedrock

1. Go to Bedrock console > Model access.
2. Click Enable specific models.
3. Select and enable the following models:
 - Titan Text Embeddings V2
 - Llama 3.3 70B Instruct
4. Click Next, accept terms, and Submit.

Edit model access

Base models (3/18)

Not seeing a model you're interested in? Check out all supported models by region [here](#).

Find model

Group by provider

	Models	Access status	Modality	EULA 
☐	▼ Amazon (4)	0/4 access granted		
☒	Titan Text Embeddings V2	 Available to request	Embedding	EULA
☐	Nova Pro Cross-region inference	 Available to request	Text & Vision	EULA
☐	Nova Lite Cross-region inference	 Available to request	Text & Vision	EULA
☐	Nova Micro Cross-region inference	 Available to request	Text	EULA
☐	▼ Anthropic (5)	0/5 access granted		
☐	Claude 3.7 Sonnet Cross-region inference	 Available to request	Text & Vision	EULA
☐	Claude 3.5 Haiku Cross-region inference	 Available to request	Text	EULA
☐	Claude 3.5 Sonnet v2 Cross-region inference	 Available to request	Text & Vision	EULA

Model	Access	Capabilities	License
☐ Claude 3.5 Sonnet v2 Cross-region inference	 Available to request	Text & Vision	EULA
☐ Claude 3.5 Sonnet Cross-region inference	 Available to request	Text & Vision	EULA
☐ Claude 3 Haiku Cross-region inference	 Available to request	Text & Vision	EULA
☐ ▼ DeepSeek (1)	0/1 access granted		
☐ DeepSeek-R1 Cross-region inference	 Available to request	Text	EULA
☐ ▼ Meta (8)	0/8 access granted		
☒ Llama 3.3 70B Instruct	 Available to request	Text	EULA
☐ Llama 3.2 1B Instruct Cross-region inference	 Available to request	Text	EULA
☐ Llama 3.2 3B Instruct Cross-region inference	 Available to request	Text	EULA
☐ Llama 3.2 11B Vision Instruct Cross-region inference	 Available to request	Text & Vision	EULA
☐ Llama 3.2 90B Vision Instruct Cross-region inference	 Available to request	Text & Vision	EULA
☐ Llama 3.1 405B Instruct Cross-region inference	 Available to request	Text	EULA
☐ Llama 3.1 70B Instruct Cross-region inference	 Available to request	Text	EULA
☒ Llama 3.1 8B Instruct Cross-region inference	 Available to request	Text	EULA

Cancel

Next

- Step 1 Edit model access
- Step 2 Add use case details for Anthropic
- Step 3 **Review and submit**

Review and submit

Step 1: Edit model access

[Edit](#)

Model access modifications (3)

Models	Modifications
Llama 3.3 70B Instruct	Request access
Titan Text Embeddings V2	Request access
Llama 3.1 8B Instruct	Request access

Step 2: Add use case details for Anthropic

Use case details

No use case details are required because no Anthropic models were requested.

 Terms

By selecting Submit, you are requesting access to the selected third party models through the AWS Marketplace. By doing so, you agree to the seller's pricing terms and End User License Agreements (EULA), and the [Bedrock Service Terms](#). You also agree and acknowledge that AWS may share information about this

Step 3
Review and submit

Model access modifications (3)

Models	Modifications
Llama 3.3 70B Instruct	Request access
Titan Text Embeddings V2	Request access
Llama 3.1 8B Instruct	Request access

Step 2: Add use case details for Anthropic

Use case details

No use case details are required because no Anthropic models were requested.

Terms

By selecting Submit, you are requesting access to the selected third party models through the AWS Marketplace. By doing so, you agree to the seller's pricing terms and End User License Agreements (EULA), and the [Bedrock Service Terms](#). You also agree and acknowledge that AWS may share information about this transaction with the respective sellers, in accordance with the [AWS Privacy Notice](#).

AWS will issue invoices and collect payments from you on behalf of the seller through your AWS account. Your use of AWS services is subject to the [AWS Customer Agreement](#) or other agreements with AWS governing your use of such services.

Cancel

Previous

Submit

To use Bedrock serverless models, account users with the correct [IAM Permissions](#) must enable access to available Bedrock foundation models (FMs). View all [Bedrock Model Terms](#) for Bedrock FMs.

Modify model access

Visit [Amazon Bedrock Quotas](#) for a quick guide to the default quotas and limits that apply to Amazon Bedrock.

Base models (18)

Collapse all

Not seeing a model you're interested in? Check out all supported models by region [here](#).

Models	Access status	Modality	EULA
▼ Amazon (4)	1/4 access granted		
Titan Text Embeddings V2	✔ Access granted	Embedding	EULA
Nova Pro Cross-region inference	⌚ Available to request	Text & Vision	EULA
Nova Lite Cross-region inference	⌚ Available to request	Text & Vision	EULA
Nova Micro Cross-region inference	⌚ Available to request	Text	EULA
▼ Anthropic (5)	0/5 access granted		
Claude 3.7 Sonnet Cross-region inference	⌚ Available to request	Text & Vision	EULA
Claude 3.5 Haiku Cross-region inference	⌚ Available to request	Text	EULA

Claude 3.7 Sonnet	Cross-region inference	 Available to request	Text & Vision	EULA
Claude 3.5 Haiku	Cross-region inference	 Available to request	Text	EULA
Claude 3.5 Sonnet v2	Cross-region inference	 Available to request	Text & Vision	EULA
Claude 3.5 Sonnet	Cross-region inference	 Available to request	Text & Vision	EULA
Claude 3 Haiku	Cross-region inference	 Available to request	Text & Vision	EULA
▼ DeepSeek (1)		0/1 access granted		
DeepSeek-R1	Cross-region inference	 Available to request	Text	EULA
▼ Meta (8)		2/8 access granted		
Llama 3.3 70B Instruct		 Access granted	Text	EULA
Llama 3.2 1B Instruct	Cross-region inference	 Available to request	Text	EULA
Llama 3.2 3B Instruct	Cross-region inference	 Available to request	Text	EULA
Llama 3.2 11B Vision Instruct	Cross-region inference	 Available to request	Text & Vision	EULA
Llama 3.2 90B Vision Instruct	Cross-region inference	 Available to request	Text & Vision	EULA
Llama 3.1 405B Instruct	Cross-region inference	 Available to request	Text	EULA
Llama 3.1 70B Instruct	Cross-region inference	 Available to request	Text	EULA
Llama 3.1 8B Instruct	Cross-region inference	 Access granted	Text	EULA

Amazon Bedrock

- ▼ Getting started
 - Overview
 - Providers
- ▼ Foundation models
 - Model catalog [New](#)
 - Marketplace deployments [New](#)
- ▼ Playgrounds
 - Chat / Text
- ▼ Builder tools
 - Agents
 - Flows
 - Knowledge Bases
 - Prompt Management
- ▼ Safeguards
 - Guardrails
- ▼ Inference and Assessment
 - Batch inference
 - Cross-region inference
 - Evaluations

Log Deliveries
Configure log deliveries and event logs in the [Edit](#) page.

Retrieval-Augmented Generation (RAG)
type
Vector store

Data source (1)

Sync

Stop sync

Add

Add documents from S3

Data sources contain information returned when querying a Knowledge Base.

☒	Data so...	Status	Data sour...	Account ID	Source Link
☒	AWS-rag-...	Syncing	S3	59379303...	s3://aws-r...

Tags
0
A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

Test Knowledge Base

Generate responses


Llama 3.1 8B In... v1

- US Meta Llama 3.1 8B Instruct

The system is syncing your data source. Wait for the sync to complete before starting next sync job.

Go to data sources

Test your Knowledge Base by running a query to generate responses. To disable response generation and only see retrieved information stored from your vector store, turn off Generate responses above.

Step 5: Sync Knowledge Base

1. In the Bedrock console, open your Knowledge Base.
2. Go to Data Sources tab.
3. Select your data source and click Sync.

The screenshot shows the Amazon Bedrock console interface. On the left is a navigation sidebar with sections: Getting started, Foundation models, Playgrounds, Builder tools, Safeguards, and Inference and Assessment. The main content area is titled 'Amazon Bedrock' and shows the 'Data Sources' tab for a Knowledge Base. It includes a 'Log Deliveries' section, a 'Retrieval-Augmented Generation (RAG)' type section, and a 'Data source (1)' section with buttons for 'Sync', 'Stop sync', 'Add', and 'Add documents from S3'. Below this is a table of data sources with columns: Data source, Status, Data source..., Account ID, and Source Link. The table shows one entry: 'AWS-rag-...' with status 'Syncing' and source 's3://aws-r-...'. A 'Tags' section is also visible at the bottom.

The screenshot shows the 'Select model' dialog in the Amazon Bedrock console. It has a search bar at the top and three main sections: 1. Categories, 2. Models, and 3. Inference. In the 'Categories' section, 'Meta' is selected under 'Model providers'. In the 'Models' section, 'Llama 3.1 8B Instruct v1' is selected under 'Models with access (2)'. In the 'Inference' section, 'US Meta Llama 3.1 8B Instruct' is selected under 'Cross-region'. At the bottom, there is a footer with a link to 'See all models here' and buttons for 'Cancel' and 'Apply'.

Embeddings model

Model

Titan Text Embeddings v2

Embeddings type

Float vector embeddings

Vector dimensions

1024

Vector database

Vector database

Amazon OpenSearch Serverless

Collection ARN

 arn:aws:aoss:us-east-2:593793030302:collection/xut9u3k0letqvn1h9368

Vector index name

bedrock-knowledge-base-default-index

Vector field name

bedrock-knowledge-base-default-vector

Text field name

AMAZON_BEDROCK_TEXT

Metadata field name

AMAZON_BEDROCK_METADATA

Multimodal storage destination

S3 URI

Test Knowledge Base >

Generate responses

 **Llama 3.3 70B I...** v1 ⓘ ✎

On-demand

 hello

 Sorry, I am unable to assist you with this request.

Enter your message here

▶ Run

infrastructure. It is not directly related to insurance, but rather a framework for organizations to manage and reduce their cyber risks.^[1]

[Show details >](#)



Explain AWS shared responsibility model as described in the guide.



The AWS Shared Responsibility Model explains that AWS is responsible for the security of the cloud, including the underlying infrastructure, while customers are responsible for security in the cloud, encompassing the management of workloads, guest operating systems, and configurations of network, host, IAM, and storage resources.^{[1][2]}

[Show details >](#)

Enter your message here

Run

Step 6: Test Your Chatbot via CLI (CloudShell)

1. Open AWS CloudShell.
2. Run the following command to test chatbot:

```
...  
aws bedrock-agent-runtime retrieve-and-generate --region us-east-2 --input '{"text":  
"What is the AWS Well-Architected Framework?"}' --retrieve-and-generate-configuration  
{  
  "knowledgeBaseConfiguration": {  
    "knowledgeBaseId": "REPLACE_WITH_YOUR_KB_ID",  
    "modelArn": "arn:aws:bedrock:us-east-2::foundation-model/meta.llama3-3-70b-  
instruct-v1:0"  
  },  
  "type": "KNOWLEDGE_BASE"  
}' --query 'output.text' --output text  
...
```

3. Replace "REPLACE_WITH_YOUR_KB_ID" with your actual Knowledge Base ID.



The screenshot shows the AWS CloudShell interface. The terminal window is titled 'CloudShell' and shows the command being executed. The output displays the AWS Shared Responsibility Model and the six pillars of the AWS Well-Architected Framework.

```
CloudShell  
us-east-2 +  
> --region us-east-2 \  
> --input '{"text": "\$QUESTION"}' \  
> --retrieve-and-generate-configuration '{  
>   "knowledgeBaseConfiguration": {  
>     "knowledgeBaseId": "ZUOKTHUJY",  
>     "modelArn": "arn:aws:bedrock:us-east-2::foundation-model/meta.llama3-3-70b-instruct-v1:0"  
>   },  
>   "type": "KNOWLEDGE_BASE"  
> }' \  
> --query 'output.text' \  
> --output text  
> }  
~ $ ask bedrock "What is the AWS Shared Responsibility Model?"  
The AWS Shared Responsibility Model is a framework that defines the responsibilities of AWS and its customers in securing solutions in the AWS cloud. AWS is responsible for the security of the cloud, which includes safe  
guarding the infrastructure that powers all services provided in the AWS cloud, such as hardware, software, networking, and facilities. On the other hand, customers are accountable for security in the cloud, which encom  
passes the management of workloads deployed in the cloud, guest operating systems, and configurations of network, host, IAM, and storage resources.  
~ $ ask bedrock "List the six pillars of the AWS Well-Architected Framework."  
The six pillars of the AWS Well-Architected Framework are operational excellence, security, reliability, performance efficiency, cost optimization, and sustainability.  
~ $
```